

---

# Shibboleth: Probe-and-Replay Model-Response Watermarking for Diffusion Language Models

---

Anonymous Author(s)

Affiliation

Address

email

## Abstract

1       Text watermarks usually hide a secret in token choices and later test those tokens  
2       without rerunning the generator. We ask whether a masked diffusion language  
3       model (DLM) can instead help verify its own output. Hiding selected parts of the  
4       context changes the probability of an emitted token in a reproducible, checkpoint-  
5       specific direction. Shibboleth uses that model response as the watermark carrier.  
6       The method first finds positions where plausible candidates exhibit both response  
7       directions. Only after a position is fixed does the key choose which direction to  
8       favour. The generator then resamples from its current distribution without changing  
9       the decoding schedule. To verify a finished document, Shibboleth reconstructs  
10      the same masked contexts and counts how often the responses agree with the  
11      key. Position selection cannot observe the requested direction, so text produced  
12      independently of the key gives an exact binomial count.  
13      At 1% FPR, Shibboleth reaches 0.80 TPR on LLaDA-8B-Instruct and 0.90 on  
14      Dream-7B-Instruct. The strongest token-local baseline reaches 0.93 on LLaDA  
15      but 0.57 or less on Dream when its LLaDA setting is transferred unchanged. At  
16      Shibboleth’s main setting, the PPL ratio is  $0.74\times$  on LLaDA and  $1.09\times$  on Dream,  
17      and changes across the measured task scores are at most 0.06. This checkpoint  
18      binding has a cost: verification needs the same model and 144 full-sequence calls  
19      per 512-token document, and dispersed edits reduce detection power.

## 20   1 Introduction

21   When language-model text resembles human writing, its origin can be difficult to establish in settings  
22   such as disinformation analysis, phishing investigation, and academic-integrity review [28]. A  
23   post-hoc detector learns recurring output patterns; a watermark instead changes generation under  
24   a secret key and lets the key holder test a later document [21, 33]. We focus on watermarking  
25   because it can offer a stated finite-sample false-positive rate, rather than relying on a classifier whose  
26   behaviour may change with the decoder [10, 19, 23, 44].

27   Nearly all language-model watermarks place the secret in token identities: the key biases a vocabulary  
28   subset or couples randomness to each draw. Their detector therefore needs only the text and key  
29   (Figure 1a) [21, 24]. This design is inexpensive, but it must coexist with the decoder. For masked  
30   DLMs, that decoder may resolve many positions in parallel and choose their order from confidence,  
31   so forcing a resolved prefix or steering the order can change the system being watermarked [17].

32   This tension suggests a different question: can the originating model, rather than the token alone,  
33   carry the signal? A masked DLM can answer a reproducible question about its own output. Take a  
34   completed sentence, hide one subset of its earlier context, and record the probability of an emitted  
35   token; then hide a second subset and record it again. Some candidates become more likely and others  
36   less likely. The sign of that change is a *model response*. Because the verifier can recreate both hidings




Figure 1: **Two places to encode a watermark.** Token-local methods alter sampled identities and verify from  $y, K$  alone. Shibboleth instead encodes which context response the checkpoint favours, so verification reruns the same  $\theta$ .

from the finished text, this response can carry a keyed signal even when the token identity alone does not (Figure 1b).

Two obstacles separate this observation from a valid detector. First, verification must recreate the generation-time state even though the full continuation is now visible. Second, the method may seek positions where both response directions are plausible, but it must finish that search before the key reveals which direction counts as a match. Reversing this order would bias the selected positions toward the requested answer.

Shibboleth addresses both obstacles. For each active block, it re-masks that block and its suffix, compares pairs of hidden-prefix probes, and selects positions whose smaller directional probability total exceeds a fixed cutoff. The key then supplies the requested direction. At a selected position, Shibboleth draws up to  $R$  candidates from the decoder’s current distribution and commits the first match; all other positions follow the reference decoder. Verification reconstructs the probes, repeats the selection rule, and evaluates an exact binomial tail. The block partition, confidence order, and checkpoint probabilities remain those of the reference decoder.

Our contributions are:

- We introduce model-response watermarking: the signal is a checkpoint’s reproducible reaction to hidden context rather than a keyed pattern in token identities.
- We give a replayable construction whose carrier positions are selected before the keyed orientation is consulted, yielding finite-sample false-positive control while preserving the reference decoding schedule (Sections 3.2–3.5).
- We evaluate two masked DLMs. At 1% FPR, Shibboleth reaches 0.80 TPR on LLaDA and 0.90 on Dream. Verification requires 144 full-sequence model calls, and detection power decreases under dispersed edits (Section 4).

## 2 Related Work

**Masked diffusion language models.** Masked DLMs generate text by repeatedly predicting unresolved positions and committing a subset of them. This extends discrete diffusion to text [3, 16, 29, 32, 37, 38] and now supports instruction-following models including LLaDA and Dream [25, 31, 42, 45]. Which positions commit, and in what order, affects both quality and speed [2, 6, 20, 26, 39, 40]. We therefore treat the deployed schedule as a constraint: Shibboleth observes the model’s masked-token predictions but does not choose a new commit order.

**Watermarks encoded in tokens.** Autoregressive watermarks usually bias a keyed vocabulary subset [21] or couple each draw to keyed randomness [1, 13, 24]. Their detectors are cheap because a key–token calculation needs no model call, and later work characterises when those detectors are reliable or undetectable [10, 22, 44]. DLM methods adapt this idea by removing prefix dependence, averaging a red–green bias over unresolved context, sampling with keyed randomness, or encoding information in the denoising order [4, 9, 14, 17, 34, 41]. Our closest baseline, dgMARK, changes which mask resolves next and later reads that order from the text. Shibboleth instead keeps the order fixed and asks the checkpoint to reproduce its response.

**Detectors that use a model.** Model access can support richer tests than a key–token hash. Likelihood-ratio detectors compare watermarked and base conditionals [18]; access to the model can improve a Neyman–Pearson score over a key-only Gumbel statistic [12]; GaussMark verifies a keyed weight perturbation with a forward–backward pass [7]; and model-theft watermarks query a suspect system [43]. Outside watermarking, DetectGPT scores text using log probabilities from the model of interest and perturbed versions of the passage; replacing the source model with a surrogate weakens detection [30]. These model-assisted detectors are most natural when a provider or central verifier can afford model access. Shibboleth shares that operational setting but adds a keyed generation channel and a document-level test. It leaves the checkpoint unchanged and keys the contexts presented to it, so one deployed model can serve many keys.

**Finite-sample detection and editing.** Watermark tests increasingly replace asymptotic  $z$ -scores with exact tails, permutation tests, or corrected maxima over windows [12, 22, 24]. After editing, a statistic that concentrates on surviving regions can outperform a document-wide sum [27]. Shibboleth uses an exact binomial tail for the unedited detector and compares document-wide, block-local, and windowed aggregation under editing in Section 4.2.

### 3 Shibboleth: Watermarking by Model Response

#### 3.1 Masked Decoding

A masked DLM predicts many unresolved positions at once rather than exposing a unique “next token.” Given a sequence containing masks, it returns a distribution at every unresolved position. Let  $x$  be the prompt,  $y = (y_1, \dots, y_N)$  the continuation, and  $I$  the positions already revealed. At any masked position  $i \notin I$ , the model provides  $p_\theta(y_i | y_I, x)$ . One forward pass over the full sequence provides these distributions for every unresolved position; we count this as one full-sequence model call.

We study the blockwise schedule used by LLaDA [2, 31]. Blocks finish from left to right, while a confidence rule decides which masks inside the active block commit at each step. We write  $B_b$  for the active block and  $F_b$  for the completed prefix before it. Shibboleth keeps both the block order and this confidence rule.

#### 3.2 Problem Setup

We target keyed provenance detection: distinguish text generated with a secret document key from text produced independently of that key, at a false-positive rate fixed in advance. The watermark may change which token is sampled, but should not replace the model’s decoder or its schedule.

**What the two parties retain.** The embedder has the model, prompt, secret key  $K$ , and target bits. The verifier later has the completed sequence, the same checkpoint,  $K$ , and public parameters—but no generation trace, stored logits, random seed, or record of rejected samples. A nonce derives a fresh document key  $K_d = \text{HMAC}(K, \text{nonce}_d)$  [5]. Separate hash domains use  $K_d$  for probe construction, carrier assignment, and the requested response directions.

**Guarantee and scope.** For any text produced independently of  $K_d$ , including output from another sampler, a detector run at level  $\alpha$  fires with probability at most  $\alpha$ . This guarantee is finite-sample and permits the method to choose a data-dependent number of carriers. It does not promise detection after editing, key disclosure, adaptive detector queries [19], or replay with a numerically different checkpoint.

**Why write during generation?** Editing only after completion offers too little freedom at most positions. Across 36 post-hoc substitution settings, the median selected position has only one candidate within a 3-nat likelihood budget; the best setting below a  $1.35 \times$  PPL budget reaches 0.10 TPR at 1% FPR. Shibboleth therefore writes while the model’s live distribution is still available.

The design follows three requirements:

1. **Replayability.** From the completed sequence, the verifier must reconstruct the same masked context that the embedder used.




Figure 2: **Probe pair and replay.** Two hidden-prefix probes split candidate tokens by which probe assigns more probability. Both groups must be viable before the key chooses one; verification recreates the same probes from the finished text.

- 123 2. **Blind selection.** Carrier positions must be fixed before the key reveals which response direction
- 124 counts as a match.
- 125 3. **Schedule preservation.** Blocks and within-block commit order must remain those of the refer-
- 126 ence decoder.

127 Replayability makes a generation-time response measurable later. Blind selection keeps adaptive

128 carrier discovery from biasing the detector, while schedule preservation limits changes to the deployed

129 generator. The next three subsections implement these requirements.

### 130 3.3 Probing the Model and Selecting Carriers

131 Figure 2 first locates the main operations within one generation-and-verification block. For each

132 unfinished block, Shibboleth asks the model the same question under several masked-prefix probes:

133 how likely is each candidate token when this part of the earlier context is missing? Comparing two

134 probes splits candidate tokens by which probe assigns them more probability. The position can carry

135 a bit only when the live decoder assigns sufficient probability to both groups; only then does the key

136 choose which group the emitted token should favour.

137 Formally, Shibboleth constructs  $L$  pseudorandom probe patterns  $D_{b,\ell} \subset F_b$ , where each pattern

138 specifies prefix positions to hide. It also masks the current block  $B_b$  and every later position. Let

139  $\lambda_{i,\ell}(v)$  denote the resulting log-probability of vocabulary item  $v$  at position  $i$  under pattern  $\ell$ . The

140 block and suffix are hidden because neither was available when the probe was created. The verifier can

141 therefore reconstruct the same state from the completed sequence instead of seeing later tokens that

142 would alter the probabilities. Probing a block costs  $L + 1$  full-sequence model calls:  $L$  masked-prefix

143 probes and one evaluation with the full prefix visible.

144 For a pair  $(\ell, \ell')$ , define the response  $g_i^{\ell, \ell'}(z) = \lambda_{i, \ell'}(z) - \lambda_{i, \ell}(z)$ . Let  $p_i^{\text{base}}$  be the block-masked

145 distribution with no hidden prefix. The aggregate probability assigned to either response direction,

146 and a symmetric score for the pair, are

$$q_{i,\pm}^{\ell, \ell'} = \sum_z p_i^{\text{base}}(z) \mathbf{1}[\pm g_i^{\ell, \ell'}(z) > 0], \quad S_i^{\ell, \ell'} = 2 \min(q_{i,+}^{\ell, \ell'}, q_{i,-}^{\ell, \ell'}). \quad (1)$$

147 The score is twice the smaller total. It approaches one when the live distribution is divided evenly

148 between the two response groups and approaches zero when one group has little probability. At

149 each position, Shibboleth keeps the pair with the largest  $S_i$  and selects the position when  $S_i > s_{\min}$ ,

150 subject to an optional per-block cap. Choosing the best pair raises the measured mean score from

151 0.28 for a fixed near/far pair to 0.45. Figure 3 shows where this criterion retains or skips positions in

152 one saved LLaDA output.

153 Crucially, swapping  $\ell$  and  $\ell'$  leaves  $S_i$  unchanged while reversing every response sign. Carrier

154 selection therefore cannot favour the direction later requested by the key.

155 After the carrier is fixed, the key assigns one group to the target bit. Formally, a separate hash domain

156 supplies  $\epsilon_i \in \{-1, +1\}$ , while  $w_i \in \{-1, +1\}$  represents the provenance or payload bit. Together

157 they decide which response direction counts as a match:

$$m_i(z) = \mathbf{1}[w_i \epsilon_i g_i(z) > 0]. \quad (2)$$




Figure 3: **Carrier availability in one generated document.** Each cell is one token position. Darker blue indicates greater balance between the response groups in Figure 2; orange outlines mark positions below the cutoff, where no bit is written.

158 If  $g_i(z) = 0$ , a second domain-separated bit assigns the outcome. Such a tie cannot be steered, but it  
 159 remains an independent fair outcome for detection.

### 160 3.4 Embedding Without Changing the Schedule

161 The reference decoder still chooses the next position and supplies its live distribution  $p_i$ . Unselected  
 162 positions are sampled once as usual. At a selected carrier, Shibboleth accepts a candidate only if it  
 163 has the requested response and is not too far below the most likely token:

$$A_i = \left\{ z : w_i \epsilon_i g_i(z) > 0 \text{ and } p_i(z) \geq \kappa \max_v p_i(v) \right\}. \quad (3)$$

164 It draws from the same  $p_i$  up to  $R$  times and commits the first member of  $A_i$ . If every proposal  
 165 misses, it takes one fresh unchecked draw. All retries reuse  $p_i$ , so they add no model call and do not  
 166 alter the commit order.

167 Let  $a_i = p_i(A_i)$  be the probability of satisfying both conditions, and let  $q_i = \sum_z p_i(z) m_i(z)$  be the  
 168 chance that an unchecked draw scores as a match. The committed token comes from  $p_i(\cdot \mid A_i)$  with  
 169 probability  $1 - (1 - a_i)^R$  and from  $p_i$  otherwise. Its match probability is therefore

$$\Pr[m_i = 1] = 1 - (1 - a_i)^R (1 - q_i). \quad (4)$$

170 With no plausibility floor and no ties,  $a_i = q_i$ , giving  $1 - (1 - q_i)^{R+1}$ . Equation (4) separates the two  
 171 controls:  $R$  determines how many chances the sampler has to find an acceptable token, while  $\kappa$  limits  
 172 how unlikely that token may be relative to the live argmax. Larger  $R$  can strengthen detection but  
 173 changes the sampling distribution;  $\kappa$  does not constrain the unchecked fallback or guarantee semantic  
 174 quality, which we measure separately.

### 175 3.5 Watermark Detection

176 The verifier re-masks each completed block and suffix, reconstructs the probe pairs and carriers, and  
 177 evaluates  $m_i(y_i)$ . It needs no rejected samples or generation trace. For provenance detection, the  
 178 selected carriers form a pool  $P$  with  $w_i = +1$ ; a payload implementation would keep a separate pool  
 179 for each message bit. The document statistic and its one-sided  $p$ -value are

$$T = \sum_{i \in P} m_i, \quad p_{\text{det}} = \Pr[\text{Binomial}(|P|, \frac{1}{2}) \geq T]. \quad (5)$$

180 Figure 4 follows a five-carrier example from replayed outcomes to the exact upper-tail probability.




Figure 4: **Document-level detection.** The verifier replays each carrier and records a match. Here four of five match; the blue bars give the exact chance of observing at least four matches under fair coins.

| Method                   | Setting                 | PPL ratio ↓   | TPR at FPR ↑ |             | Full-sequence model calls ↓ |              |
|--------------------------|-------------------------|---------------|--------------|-------------|-----------------------------|--------------|
|                          |                         |               | 1%           | 0.1%        | Generation                  | Verification |
| <i>LLaDA-8B-Instruct</i> |                         |               |              |             |                             |              |
| KGW [21]                 | $\delta = 1$            | 1.22×         | 0.90         | 0.83        | 512                         | <b>0</b>     |
| dgMARK [17]              | $k = 1$                 | 1.45×         | 0.83         | 0.80        | 512                         | <b>0</b>     |
| dgMARK [17]              | +3-beam                 | 1.21×         | <b>0.93</b>  | <b>0.93</b> | ~2048                       | <b>0</b>     |
| Shibboleth               | $R = 1$                 | 0.97×         | 0.43         | 0.20        | 656                         | 144          |
| Shibboleth               | $R = 8, \kappa = 0.1$   | 0.87×         | 0.70         | 0.60        | 656                         | 144          |
| Shibboleth               | $R = 16, \kappa = 0.05$ | <b>0.74</b> × | 0.80         | 0.60        | 656                         | 144          |
| <i>Dream-7B-Instruct</i> |                         |               |              |             |                             |              |
| KGW [21]                 | $\delta = 1$            | 1.84×         | 0.57         | 0.37        | 512                         | <b>0</b>     |
| dgMARK [17]              | $k = 1$                 | 1.75×         | 0.37         | 0.27        | 512                         | <b>0</b>     |
| dgMARK [17]              | +3-beam                 | <b>0.71</b> × | 0.30         | 0.30        | ~2048                       | <b>0</b>     |
| Shibboleth               | $R = 1$                 | 1.34×         | 0.40         | 0.27        | 656                         | 144          |
| Shibboleth               | $R = 8, \kappa = 0.1$   | 0.79×         | 0.70         | 0.67        | 656                         | 144          |
| Shibboleth               | $R = 16, \kappa = 0.05$ | 1.09×         | <b>0.90</b>  | <b>0.90</b> | 656                         | 144          |

Table 1: **Detectability and cost at 512 tokens.** PPL ratio is measured against each method’s own unwatermarked run and should be interpreted within that method. Shibboleth uses  $n = 30$  documents ( $n = 20$  for  $R = 16$ ); Dream PPL uses 8–13 successfully scored document pairs per setting. Settings chosen on LLaDA are transferred to Dream without retuning. One model call is one forward pass over the current full sequence; these counts measure algorithmic work, not runtime.  $\delta$  is KGW’s logit bias and  $k$  is dgMARK’s candidate count.

**Proposition 1** (Exact false-positive control). Assume the hash domains behave as independent random functions. For any text written without the document key, condition on the selected carriers, probe pairs, responses, and target bits. The orientation bits remain independent fair coins. Hence the  $m_i$  are independent Bernoulli( $\frac{1}{2}$ ) variables,  $T \sim \text{Binomial}(|P|, \frac{1}{2})$ , and thresholding Eq. (5) at  $\alpha$  gives false-positive rate at most  $\alpha$  for every document length and carrier count.

**Why adaptive carrier selection is allowed.** The text and model may determine the number of carriers, the probe pair, and the response size. What matters is that each choice is fixed before the independent orientation bit is used. Flipping  $\epsilon_i$  then flips the match whenever  $g_i(y_i) \neq 0$ ; the tie rule supplies a fresh fair bit otherwise. The detector therefore needs neither an asymptotic approximation nor a calibration corpus. The guarantee still assumes text produced independently of the key and domain separation among the key’s several roles.

## 4 Experiments

The evaluation asks four questions: how strongly the watermark can be detected, whether one setting transfers across checkpoints, what it costs in computation and output quality, and how editing changes detection power.

**Models and data.** We evaluate LLaDA-8B-Instruct [31] and Dream-7B-Instruct [42] in FP16 on one TITAN RTX. Detectability uses C4 realnewslike prompts [36], truncated to 300 characters, and 512-token continuations divided into 32-token blocks. We correct Dream’s one-position logit offset in its wrapper, then use the same position convention for generation and replay.

**Decoders and settings.** Shibboleth and dgMARK use each model’s confidence schedule with multinomial sampling at temperature 0.8; dgMARK additionally uses top- $k = 3$ . KGW requires strictly left-to-right decoding. Shibboleth uses  $L = 8$  probe patterns, threshold  $s_{\min} = 0.5$ , retry budget  $R$ , and optional plausibility floor  $\kappa$ . We evaluate provenance ( $w_i = +1$ ), not payload recovery. Here one model call means one forward pass over the current full sequence. With 16 blocks and  $L + 1 = 9$  probe evaluations, verification replays  $16 \times 9 = 144$  calls; generation uses 512 ordinary decoding calls plus the same 144 probe calls, for 656 in total.

**Metrics.** Detection is TPR at the analytic FPR thresholds in Eq. (5). Perplexity is scored by GPT-2-large [35] as  $\exp(\text{mean}(\text{NLL}_{\text{wm}} - \text{NLL}_{\text{unwm}}))$ . A PPL ratio therefore compares a method with its



| Method                   | MMLU Acc $\uparrow$       | GSM8K Acc $\uparrow$      | HumanEval Pass@1 $\uparrow$ |
|--------------------------|---------------------------|---------------------------|-----------------------------|
| <i>LLaDA-8B-Instruct</i> |                           |                           |                             |
| KGW [21]                 | 0.530 $\rightarrow$ 0.535 | 0.640 $\rightarrow$ 0.740 | 0.329 $\rightarrow$ 0.293   |
| dgMARK [17]              | 0.605 $\rightarrow$ 0.530 | 0.740 $\rightarrow$ 0.660 | 0.409 $\rightarrow$ 0.341   |
| dgMARK +3-beam           | 0.605 $\rightarrow$ 0.530 | 0.740 $\rightarrow$ 0.660 | 0.409 $\rightarrow$ 0.280   |
| Shibboleth               | 0.595 $\rightarrow$ 0.575 | 0.640 $\rightarrow$ 0.700 | 0.335 $\rightarrow$ 0.384   |
| <i>Dream-7B-Instruct</i> |                           |                           |                             |
| KGW                      | 0.640 $\rightarrow$ 0.625 | 0.700 $\rightarrow$ 0.700 | —                           |
| dgMARK                   | 0.740 $\rightarrow$ 0.570 | 0.840 $\rightarrow$ 0.560 | —                           |
| dgMARK +3-beam           | 0.740 $\rightarrow$ 0.645 | 0.840 $\rightarrow$ 0.560 | —                           |
| Shibboleth               | 0.695 $\rightarrow$ 0.710 | 0.660 $\rightarrow$ 0.680 | —                           |

Table 2: **Task accuracy before and after watermarking.** Each cell is unwatermarked  $\rightarrow$  watermarked under that method’s decoder. MMLU uses 200 questions, GSM8K 50 problems, and HumanEval 164 executed tasks. Approximate standard errors are 0.035, 0.07, and 0.037, respectively; — denotes an unmeasured pair. Shibboleth uses  $R = 16$ ,  $\kappa = 0.05$ .

own unwatermarked decoder; it is not a semantic-quality score and is only indicative across different decoders. The baselines are KGW [21] and dgMARK [17].

#### 4.1 Main Results and Analysis

Table 1 shows the tradeoff: Shibboleth transfers strongly to Dream but replay is costly, while two cheaper baselines detect LLaDA output more often.

**Retries trade sampling distortion for detection.** Table 1 gives the 512-token detection, PPL, and model-call comparison. On LLaDA,  $R = 1$  reaches 0.43 TPR at 1% FPR. On 30 held-out prompts,  $R = 2, 4, 8$  increase TPR to 0.73, 0.83, 0.87, consistent with Eq. (4). Retries add no model evaluation, but without a floor the PPL ratio rises from  $1.14\times$  at  $R = 2$  to  $1.61\times$  at  $R = 8$ . The floor removes low-probability guided draws:  $(R, \kappa) = (8, 0.1)$  gives 0.70 TPR at  $0.87\times$  PPL, and  $(16, 0.05)$  gives 0.80 at  $0.74\times$ . Ratios below one mean only that GPT-2-large prefers the retained high-probability tokens.

**The transferred setting behaves differently across checkpoints.** We carry every LLaDA setting to Dream without retuning. KGW reaches 0.57 TPR at 1% FPR and produces 84% duplicate bigrams under forced left-to-right decoding. Direct and three-beam dgMARK reach 0.37 and 0.30 TPR. Shibboleth at  $(16, 0.05)$  reaches 0.90 at both reported FPR levels. Its PPL ratio is  $0.79\times$  at  $(8, 0.1)$  and  $1.09\times$  at  $(16, 0.05)$ , based on 8–13 scored pairs per setting. Thus detectability transfers in this experiment, whereas the setting with the lowest LLaDA PPL ratio does not retain that ratio on Dream.

**Task accuracy exposes decoder-specific costs.** Table 2 shows each method’s own unwatermarked and watermarked scores side by side; Appendix Table 5 gives the full grid. Shibboleth changes by at most 0.06. On LLaDA HumanEval, dgMARK changes from 0.409 to 0.341 direct and 0.280 with three beams; on Dream, MMLU changes from 0.740 to 0.570 and GSM8K from 0.840 to 0.560. On 256-token LLaDA GSM8K answers, Shibboleth and direct dgMARK reach 0.12 TPR at 1% FPR, three-beam dgMARK 0.44, and KGW zero; 512-token rates therefore do not transfer to short answers.

**Checkpoint binding costs model calls.** On LLaDA, Shibboleth’s 0.80/0.60 TPR at 1%/0.1% FPR trails three-beam dgMARK’s 0.93/0.93 and KGW’s 0.90/0.83, while requiring 144 verification calls instead of zero. KGW’s left-to-right decoder also repeats heavily: 37% of unwatermarked documents cross the threshold before bigram deduplication, and two of thirty cross 0.1% afterward. With 20–30 documents, TPR differences near 0.1 remain imprecise. Verification cost is a dial rather than a fixed price: halving the probe patterns to  $L = 4$  cuts the replay to 80 calls and keeps 0.70 TPR at 1% FPR against 0.77 at  $L = 8$ , while  $L = 2$  saves only 32 further calls and falls to 0.57 (Appendix Table 6).

## 242 4.2 Calibration and Robustness

243 **Observed false-positive rates are below nominal levels.** Across 20 keys and 120 unwatermarked  
244 LLaDA generations per key, mean FPR is 0.0121/0.0037/0.0000 at nominal 0.10/0.05/0.01; the  
245 largest per-key rates are 0.0583/0.0250/0.0000. Every Dream cohort has zero false positives at 1%  
246 (Appendix Table 3).

247 **Dispersed edits break prefix synchronisation.** At  $R = 1$ , random 10% re-denoising or deletion  
248 reduces TPR from 0.35 to 0.05. At  $R = 16$ ,  $\kappa = 0.05$ , random 5% and 10% edits leave 0.88 and 0.50  
249 TPR, whereas a contiguous 10% or 30% span leaves 0.88 because earlier carriers survive. Blockwise  
250 Bonferroni and 128-token windowing reach only 0.38 and 0.44 under random 10% editing, below  
251 pooled counting at 0.50 (Appendix Table 4).

## 252 5 Discussion

253 **Deployment role.** Shibboleth is most useful when the question concerns a registered model release  
254 — whether that checkpoint is consistent with a disputed document. Providers, centralised services,  
255 or auditors can retain the model without changing its production decoder. The method is a poor  
256 fit for browser-scale classification, unknown source models, short fragments, or heavily edited text.  
257 An inexpensive token-local detector can handle routine cases, with replay reserved for disputes that  
258 justify model-specific computation. We do not establish transfer across checkpoint copies.

259 **Reducing the verifier burden.** Clients need not execute the model themselves. A provider can  
260 replay a document on the registered checkpoint and return a signed report; an independent auditor can  
261 instead use escrowed weights. Batching improves hardware utilisation, but each document induces  
262 different masked contexts, so replay cannot become a fixed lookup; a surrogate checkpoint or cached  
263 responses may reconstruct different carriers and would need their own false-positive analysis.

264 **Trust boundary.** A hosted verifier assumes that the service runs the registered checkpoint and  
265 receives the submitted document. A signed report records the service’s claim; it does not prove  
266 correct execution or protect document privacy. Checkpoint and container hashes, append-only audit  
267 records, and independent execution can strengthen this arrangement; public proof of correct replay,  
268 attestation, service compromise, and adversarial requests are not evaluated here.

269 **Interpreting a detection result.** A low  $p$ -value has a narrow interpretation: under the registered  
270 key and replay procedure, the document contains more checkpoint-consistent carrier responses than  
271 expected by chance. It does not identify the person who requested the text, establish that every  
272 passage came from the model, or rule out editing; those questions require separate records. A report  
273 should therefore include the  $p$ -value, carrier count, key identifier, checkpoint hash, precision, probe  
274 construction, and threshold; the count separates a short, inconclusive document from checkpoint  
275 disagreement.

## 276 6 Conclusion

277 Shibboleth encodes a keyed signal in reproducible model responses to hidden context. Selecting  
278 a carrier with a symmetric score before revealing the keyed direction yields an exact binomial test  
279 even with a data-dependent carrier count. The method preserves the deployed schedule and reaches  
280 0.80 TPR on LLaDA and 0.90 on Dream without retuning, but needs 144 same-checkpoint calls  
281 per 512-token document, so it complements cheaper token-local detectors, which perform better on  
282 LLaDA.

283 **Limitations.** The false-positive guarantee does not imply editing tolerance: random 10% re-  
284 denoising reduces TPR from 0.35 to 0.05 at  $R = 1$  and from 0.88 to 0.50 at  $R = 16$ , while a  
285 contiguous span leaves earlier carriers intact; content-anchored probes may help, though the tested  
286 block-local and windowed variants did not. A PPL ratio below one does not imply better text, and  
287 the task cohorts are too small to interpret modest changes. Replay requires the same checkpoint  
288 and precision; evaluation covers two masked DLMS, one corpus, 16–50 documents per cohort, and  
289 provenance only.



## References

- [1] Scott Aaronson and Hendrik Kirchner. Watermarking GPT outputs, 2023. <https://www.scottaaronson.com/talks/watermark.ppt>.
- [2] Marianne Arriola, Subham Sekhar Sahoo, Aaron Gokaslan, Zhihan Yang, Zhixuan Qi, Jiaqi Han, Justin T. Chiu, and Volodymyr Kuleshov. Block diffusion: Interpolating between autoregressive and diffusion language models. In *ICLR*, 2025.
- [3] Jacob Austin, Daniel D. Johnson, Jonathan Ho, Daniel Tarlow, and Rianne van den Berg. Structured denoising diffusion models in discrete state-spaces. In *NeurIPS*, 2021.
- [4] Arka Bagchi, Akhil Bhimaraju, Moulik Choraria, Daniel Alabi, and Lav R. Varshney. Watermarking discrete diffusion language models. *arXiv preprint arXiv:2511.02083*, 2025.
- [5] Mihir Bellare, Ran Canetti, and Hugo Krawczyk. Keying hash functions for message authentication. In *CRYPTO*, 1996.
- [6] Heli Ben-Hamu, Itai Gat, Daniel Severo, Niklas Nolte, and Brian Karrer. Accelerated sampling from masked diffusion models via entropy bounded unmasking. In *NeurIPS*, 2025.
- [7] Adam Block, Alexander Rakhlin, and Ayush Sekhari. GaussMark: A practical approach for structural watermarking of language models. In *ICML*, 2025.
- [8] Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde de Oliveira Pinto, Jared Kaplan, et al. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*, 2021.
- [9] Ruibo Chen, Yihan Wu, Yanshuo Chen, Chenxi Liu, Junfeng Guo, and Heng Huang. A watermark for order-agnostic language models. In *ICLR*, 2025.
- [10] Miranda Christ, Sam Gunn, and Or Zamir. Undetectable watermarks for language models. In *COLT*, 2024.
- [11] Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021.
- [12] Pierre Fernandez, Antoine Chaffin, Karim Tit, Vivien Chappelier, and Teddy Furon. Three bricks to consolidate watermarks for large language models. In *2023 IEEE International Workshop on Information Forensics and Security (WIFS)*, 2023.
- [13] Jiayi Fu, Xuandong Zhao, Ruihan Yang, Yuansen Zhang, Jiangjie Chen, and Yanghua Xiao. GumbelSoft: Diversified language model watermarking via the GumbelMax-Trick. In *ACL*, 2024.
- [14] Thibaud Gloaguen, Robin Staab, Nikola Jovanović, and Martin Vechev. Watermarking diffusion language models. In *ICLR*, 2026. arXiv:2509.24368.
- [15] Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. Measuring massive multitask language understanding. In *ICLR*, 2021.
- [16] Jonathan Ho, Ajay Jain, and Pieter Abbeel. Denoising diffusion probabilistic models. In *NeurIPS*, 2020.
- [17] Pyo Min Hong and Albert No. dgMARK: Decoding-Guided Watermarking for Diffusion Language Models. In *ICML*, 2026. arXiv:2601.22985.
- [18] Zhengmian Hu, Lichang Chen, Xidong Wu, Yihan Wu, Hongyang Zhang, and Heng Huang. Unbiased watermark for large language models. In *The Twelfth International Conference on Learning Representations (ICLR)*, 2024.
- [19] Nikola Jovanović, Robin Staab, and Martin Vechev. Watermark stealing in large language models. In *ICML*, 2024.

- [20] Jaeyeon Kim, Kulin Shah, Vasilis Kontonis, Sham M. Kakade, and Sitan Chen. Train for the worst, plan for the best: Understanding token ordering in masked diffusions. In *ICML*, 2025.
- [21] John Kirchenbauer, Jonas Geiping, Yuxin Wen, Jonathan Katz, Ian Miers, and Tom Goldstein. A watermark for large language models. In *ICML*, 2023.
- [22] John Kirchenbauer, Jonas Geiping, Yuxin Wen, Manli Shu, Khalid Saifullah, Kezhi Kong, Kasun Fernando, Aniruddha Saha, Micah Goldblum, and Tom Goldstein. On the reliability of watermarks for large language models. In *ICLR*, 2024.
- [23] Kalpesh Krishna, Yixiao Song, Marzena Karpinska, John F. Wieting, and Mohit Iyyer. Paraphrasing evades detectors of AI-generated text, but retrieval is an effective defense. In *NeurIPS*, 2023.
- [24] Rohith Kuditipudi, John Thickstun, Tatsunori Hashimoto, and Percy Liang. Robust distortion-free watermarks for language models. *TMLR*, 2024.
- [25] Inception Labs, Samar Khanna, Siddhant Kharbanda, Shufan Li, Harshit Varma, Eric Wang, Sarah Birnbaum, Ziyang Luo, Yanis Miraoui, and Akash Palrecha. Mercury: Ultra-fast language models based on diffusion. *arXiv preprint arXiv:2506.17298*, 2025.
- [26] Chia-Ming Lee, Shao-Kai Liu, Ming-Ching Chang, Xin Li, Yu-Lun Liu, and Chih-Chung Hsu. Commit locally, exit globally: Coordinating adaptive sampling and early exit in diffusion language models, 2026. URL <https://arxiv.org/abs/2607.28166>.
- [27] Xiang Li, Feng Ruan, Huiyuan Wang, Qi Long, and Weijie J. Su. Robust detection of watermarks for large language models under human edits. *Journal of the Royal Statistical Society: Series B*, 2025.
- [28] Yuqing Liang, Jiancheng Xiao, Wensheng Gan, and Philip S. Yu. Watermarking techniques for large language models: A survey. *Artificial Intelligence Review*, 2026.
- [29] Aaron Lou, Chenlin Meng, and Stefano Ermon. Discrete diffusion modeling by estimating the ratios of the data distribution. In *ICML*, 2024.
- [30] Eric Mitchell, Yoonho Lee, Alexander Khazatsky, Christopher D. Manning, and Chelsea Finn. DetectGPT: Zero-shot machine-generated text detection using probability curvature. In *Proceedings of the 40th International Conference on Machine Learning*, 2023.
- [31] Shen Nie, Fengqi Zhu, Zebin You, Xiaolu Zhang, Jingyang Ou, Jun Hu, Jun Zhou, Yankai Lin, Ji-Rong Wen, and Chongxuan Li. Large language diffusion models. In *NeurIPS*, 2025.
- [32] Jingyang Ou, Shen Nie, Kaiwen Xue, Fengqi Zhu, Jiacheng Sun, Zhenguo Li, and Chongxuan Li. Your absorbing discrete diffusion secretly models the conditional distributions of clean data. In *ICLR*, 2025.
- [33] Fabien A. P. Petitcolas, Ross J. Anderson, and Markus G. Kuhn. Information hiding — a survey. In *Proceedings of the IEEE*, 1999.
- [34] Ohad Raban, Ethan Fetaya, and Gal Chechik. LR-DWM: Efficient watermarking for diffusion language models. *arXiv preprint arXiv:2601.12376*, 2026.
- [35] Alec Radford, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever. Language models are unsupervised multitask learners. Technical report, OpenAI, 2019.
- [36] Colin Raffel, Noam Shazeer, Adam Roberts, Katherine Lee, Sharan Narang, Michael Matena, Yanqi Zhou, Wei Li, and Peter J. Liu. Exploring the limits of transfer learning with a unified text-to-text transformer. *Journal of Machine Learning Research*, 21(140):1–67, 2020.
- [37] Subham Sekhar Sahoo, Marianne Arriola, Aaron Gokaslan, Edgar Marroquin Marroquin, Alexander M. Rush, Yair Schiff, Justin T. Chiu, and Volodymyr Kuleshov. Simple and effective masked diffusion language models. In *NeurIPS*, 2024.
- [38] Jiaxin Shi, Kehang Han, Zhe Wang, Arnaud Doucet, and Michalis Titsias. Simplified and generalized masked diffusion for discrete data. In *NeurIPS*, 2024.

- 383 [39] Qingyan Wei, Yaojie Zhang, Zhiyuan Liu, Pengfei Zeng, Yu Wang, Biao Qi, Dong Liu, and  
 384 Linfeng Zhang. Accelerating diffusion large language models with SlowFast sampling: The  
 385 three golden principles. In *ICLR*, 2026.
- 386 [40] Chengyue Wu, Hao Zhang, Shuchen Xue, Zhijian Liu, Shizhe Diao, Ligeng Zhu, Ping Luo,  
 387 Song Han, and Enze Xie. Fast-dLLM: Training-free acceleration of diffusion LLM by enabling  
 388 KV cache and parallel decoding. In *ICLR*, 2026.
- 389 [41] Linyu Wu, Linhao Zhong, Wenjie Qu, Yifan Li, Yue Liu, Shengfang Zhai, Chunhua Shen, and  
 390 Jiaheng Zhang. DMark: Order-agnostic watermarking for diffusion large language models.  
 391 *arXiv preprint arXiv:2510.02902*, 2025.
- 392 [42] Jiacheng Ye, Zhihui Xie, Lin Zheng, Jiahui Gao, Zirui Wu, Xin Jiang, Zhenguo Li, and Lingpeng  
 393 Kong. Dream 7B: Diffusion large language models. *arXiv preprint arXiv:2508.15487*, 2025.
- 394 [43] Xuandong Zhao, Yu-Xiang Wang, and Lei Li. Protecting language generation models via  
 395 invisible watermarking. In *ICML*, 2023.
- 396 [44] Xuandong Zhao, Prabhanjan V. Ananth, Lei Li, and Yu-Xiang Wang. Provable robust water-  
 397 marking for AI-generated text. In *ICLR*, 2024.
- 398 [45] Fengqi Zhu, Rongzhen Wang, Shen Nie, Xiaolu Zhang, Chunwei Wu, Jun Hu, Jun Zhou, Jianfei  
 399 Chen, Yankai Lin, Ji-Rong Wen, and Chongxuan Li. LLaDA 1.5: Variance-reduced preference  
 400 optimization for large language diffusion models. *arXiv preprint arXiv:2505.19223*, 2025.

## A Additional Tables

Table 3 reports the measured per-key false-positive rate of the deployed decision rule. Table 4 summarises the post-editing evaluation under the deployed pooled detector. Table 5 expands Table 2 into the full task-accuracy grid, including the greedy-decoding cells. Table 6 shows how detection and verification cost trade off as the number of probe patterns varies.

| Nominal $\alpha$ | Mean FPR | Highest per-key FPR |
|------------------|----------|---------------------|
| 0.10             | 0.0121   | 0.0583              |
| 0.05             | 0.0037   | 0.0250              |
| 0.01             | 0.0000   | 0.0000              |

Table 3: **Observed false-positive rates across keys.** The evaluation uses 20 keys and 120 unwatermarked generations per key under the deployed decision rule.

|                                                                                                            |                     | TPR at FPR $\uparrow$ |      |
|------------------------------------------------------------------------------------------------------------|---------------------|-----------------------|------|
| Attack                                                                                                     | Strength            | 1%                    | 0.1% |
| <i>Configuration A: <math>R = 1</math>, 512 tokens (<math>n = 20</math>)</i>                               |                     |                       |      |
| None                                                                                                       | —                   | 0.35                  | —    |
| Re-denoising (same model, argmax)                                                                          | 10% of positions    | 0.05                  | —    |
| Deletion                                                                                                   | 10% of positions    | 0.05                  | —    |
| Insertion                                                                                                  | 10% of positions    | —                     | —    |
| Substitution (random)                                                                                      | 10% of positions    | —                     | —    |
| Copy-paste                                                                                                 | 25% kept, 1 segment | —                     | —    |
| Paraphrase (DIPPER) [23]                                                                                   | lex 20, order 0     | —                     | —    |
| <i>Configuration B: <math>R = 16</math>, <math>\kappa = 0.05</math>, 1024 tokens (<math>n = 16</math>)</i> |                     |                       |      |
| None                                                                                                       | —                   | 0.88                  | 0.88 |
| Re-denoising (random positions)                                                                            | 5% of positions     | 0.88                  | 0.75 |
|                                                                                                            | 10% of positions    | 0.50                  | 0.38 |
| Re-denoising (contiguous span)                                                                             | 10% of positions    | 0.88                  | 0.88 |
|                                                                                                            | 30% of positions    | 0.88                  | 0.88 |

Table 4: **Detection after editing.** The pooled detector uses the same analytic thresholds as Table 1. Random edits disrupt later prefixes; a contiguous span preserves all earlier carriers. Because the cohorts are small, one document changes TPR by 0.05 when  $n = 20$  and about 0.06 when  $n = 16$ . An em dash denotes an unmeasured setting.

| Model             | Method          | Greedy decoding        |                         |                                | Multinomial decoding   |                         |                                |
|-------------------|-----------------|------------------------|-------------------------|--------------------------------|------------------------|-------------------------|--------------------------------|
|                   |                 | MMLU<br>Acc $\uparrow$ | GSM8K<br>Acc $\uparrow$ | HumanEval<br>Pass@1 $\uparrow$ | MMLU<br>Acc $\uparrow$ | GSM8K<br>Acc $\uparrow$ | HumanEval<br>Pass@1 $\uparrow$ |
| LLaDA-8B-Instruct | Non-watermarked | 0.605                  | 0.740                   | 0.409                          | 0.595                  | 0.640                   | 0.335                          |
|                   | KGW [21]        | —                      | —                       | —                              | 0.535                  | 0.740                   | 0.293                          |
|                   | dgMARK [17]     | —                      | —                       | —                              | 0.530                  | 0.660                   | 0.341                          |
|                   | dgMARK +3-beam  | —                      | —                       | —                              | 0.530                  | 0.660                   | 0.280                          |
|                   | Shibboleth      | —                      | —                       | —                              | 0.575                  | 0.700                   | 0.384                          |
| Dream-7B-Instruct | Non-watermarked | 0.740                  | 0.840                   | —                              | 0.695                  | 0.660                   | —                              |
|                   | KGW             | —                      | —                       | —                              | 0.625                  | 0.700                   | —                              |
|                   | dgMARK          | —                      | —                       | —                              | 0.570                  | 0.560                   | —                              |
|                   | dgMARK +3-beam  | —                      | —                       | —                              | 0.645                  | 0.560                   | —                              |
|                   | Shibboleth      | —                      | —                       | —                              | 0.710                  | 0.680                   | —                              |

Table 5: **Text generation quality, full grid.** Accuracy on MMLU [15] and GSM8K [11] and pass@1 on HumanEval [8] under greedy and multinomial decoding, with the unwatermarked run of each model as the reference row of its block. Shibboleth uses  $R = 16$ ,  $\kappa = 0.05$ ; 50 problems per GSM8K cell, so differences in individual scores should be interpreted cautiously. Each watermark decodes under its own regime—KGW left-to-right (its green list requires a resolved prefix; its own control scores 0.640), dgMARK under confidence decoding (its argmax control is the greedy reference cell)—so rows are read against their own controls, not across. An em dash denotes an unmeasured setting.

| $L$ | Verify calls $\downarrow$ | TPR at FPR $\uparrow$ |      |      | Sync rate |
|-----|---------------------------|-----------------------|------|------|-----------|
|     |                           | 5%                    | 1%   | 0.1% |           |
| 2   | 48                        | 0.73                  | 0.57 | 0.50 | 0.629     |
| 4   | 80                        | 0.77                  | 0.70 | 0.67 | 0.650     |
| 8   | 144                       | 0.77                  | 0.77 | 0.73 | 0.658     |

Table 6: **Fewer probe patterns make verification cheaper.** Each row halves the number of probe patterns  $L$ ; verification replays  $16(L + 1)$  model calls. Detection is measured at  $R = 16$ ,  $\kappa = 0.05$  on LLaDA, 512 tokens,  $n = 30$  fresh prompts. Going from  $L = 8$  to  $L = 4$  keeps most of the detection at a little over half the cost; going to  $L = 2$  saves little more and loses about a tenth of TPR at every operating point.