
Shibboleth: Probe-and-Replay Model-Response Watermarking for Diffusion Language Models

Anonymous Author(s)

Affiliation

Address

email

Abstract

Most text watermarks encode a secret in the chosen tokens and later inspect the text alone. We study a different channel available in masked diffusion language models (DLMs): after generation, the same model can reveal how an emitted token’s probability changes when selected parts of its context are hidden. This reproducible response ties the watermark to both the text and the checkpoint.

Shibboleth turns that response into a detector. It first finds positions whose candidate tokens span both response directions. Only after a position is selected does the key choose which direction to favour. The generator then resamples from the current model distribution, without changing the decoding schedule. Verification reconstructs the same masked contexts from the finished text and counts keyed sign agreements. Because position selection cannot observe the requested sign, text generated independently of the key produces an exact binomial count.

At 1% FPR, Shibboleth reaches 0.80 TPR on LLaDA-8B-Instruct and 0.90 on Dream-7B-Instruct. Token-based detectors reach 0.90–0.93 on LLaDA, but Shibboleth is the only evaluated row there with a PPL ratio below its own unwatermarked run; on Dream it has the highest TPR under settings transferred unchanged from LLaDA. Verification requires 144 masked evaluations per 512-token document, access to the same checkpoint, and loses power under dispersed editing.

1 Introduction

Language-model text can be difficult for readers to distinguish from human writing, creating provenance risks in disinformation, phishing, and academic-integrity review [25]. A post-hoc detector learns the output patterns of a generator; a watermark instead changes generation under a secret key and lets the key holder test a later document [18, 29]. We study watermarking because it permits a stated false-positive rate, including at finite sample size, rather than relying on a classifier whose behaviour may change with the decoder [9, 16, 20, 40].

Nearly all language-model watermarks place the secret in token identities: the key biases a vocabulary subset or couples randomness to each draw. Their detector therefore needs only the text and key (Figure 1a) [18, 21]. This design is inexpensive, but it must coexist with the decoder. For masked DLMs, that decoder may resolve many positions in parallel and choose their order from confidence, so forcing a resolved prefix or steering the order can change the system being watermarked [14].

Our starting observation is simple: a masked DLM can answer a reproducible question about its own output. Take a completed sentence, hide one subset of its earlier context, and record the probability of an emitted token; then hide a second subset and record it again. Some candidates become more likely and others less likely. The sign of that change is a *model response*. Because the verifier can recreate both hidings from the finished text, this response can carry a keyed signal even when the token identity alone does not (Figure 1b).




Figure 1: **Two places to encode a watermark.** Token-local methods alter sampled identities and verify from y, K alone. Shibboleth instead encodes which context response the checkpoint favours, so verification reruns the same θ .

Two details determine whether this idea is statistically sound. First, verification must recreate exactly the state used during generation, although later tokens are now visible. Second, a method may choose positions where both response directions are plausible, but it must do so before the key reveals which direction will count as a match. Otherwise position selection itself raises the false-positive rate.

Shibboleth addresses both points. For each active block, it re-masks that block and its suffix, compares pairs of hidden-prefix probes, and selects positions that retain probability mass on both response directions. The key then supplies the requested direction. At a selected position, Shibboleth draws up to R candidates from the decoder’s current distribution and commits the first match; all other positions follow the reference decoder. Verification reconstructs the probes, repeats the selection rule, and evaluates an exact binomial tail. The block partition, confidence order, and checkpoint probabilities remain those of the reference decoder.

Our contributions are:

- We introduce model-response watermarking, where the signal is a checkpoint’s reproducible reaction to hidden context rather than a keyed pattern in token identities.
- We give a replayable construction whose carrier positions are selected before the keyed orientation is consulted, yielding finite-sample false-positive control while preserving the reference decoding schedule (Sections 3.2–3.5).
- We evaluate two masked DLMs. At 1% FPR, Shibboleth reaches 0.80 TPR on LLaDA and 0.90 on Dream, while requiring 144 masked evaluations for verification and degrading under dispersed edits (Section 4).

2 Related Work

Masked diffusion language models. Masked DLMs generate text by repeatedly predicting unresolved positions and committing a subset of them. This extends discrete diffusion to text [3, 13, 26, 28, 33, 34] and now supports instruction-following models including LLaDA and Dream [22, 27, 38, 41]. Which positions commit, and in what order, affects both quality and speed [2, 6, 17, 23, 35, 36]. We therefore treat the deployed schedule as a constraint: Shibboleth observes the model’s masked-token predictions but does not choose a new commit order.

Watermarks encoded in tokens. Autoregressive watermarks usually bias a keyed vocabulary subset [18] or couple each draw to keyed randomness [1, 11, 21]. Their detectors are cheap because a key–token calculation needs no model call, and later work characterises when those detectors are reliable or undetectable [9, 19, 40]. DLM methods adapt this idea by removing prefix dependence, averaging a red–green bias over unresolved context, sampling with keyed randomness, or encoding information in the denoising order [4, 8, 12, 14, 30, 37]. Our closest baseline, dgMARK, changes which mask resolves next and later reads that order from the text. Shibboleth instead keeps the order fixed and asks the checkpoint to reproduce its response.

Detectors that use a model. Model access can support richer tests than a key–token hash. Likelihood-ratio detectors compare watermarked and base conditionals [15]; access to the model can improve a Neyman–Pearson score over a key-only Gumbel statistic [10]; GaussMark verifies a keyed weight perturbation with a forward–backward pass [7]; and model-theft watermarks query a suspect

77 system [39]. Shibboleth differs in where the key lives: it leaves the checkpoint unchanged and keys
 78 the contexts presented to it, so one deployed model can serve many keys.

79 **Finite-sample detection and editing.** Watermark tests increasingly replace asymptotic z -scores
 80 with exact tails, permutation tests, or corrected maxima over windows [10, 19, 21]. After editing, a
 81 statistic that concentrates on surviving regions can outperform a document-wide sum [24]. Shibboleth
 82 uses an exact binomial tail for the unedited detector and compares document-wide, block-local, and
 83 windowed aggregation under editing in Section 4.2.

84 3 Shibboleth: Watermarking by Model Response

85 3.1 Masked Decoding

86 A masked DLM does not have a unique “next token.” It sees a sequence containing masks and returns
 87 a distribution at every unresolved position. Let x be the prompt, $y = (y_1, \dots, y_N)$ the continuation,
 88 and I the positions already revealed. At any masked position $i \notin I$, the model provides $p_\theta(y_i \mid y_I, x)$.
 89 One sequence evaluation provides these distributions for all unresolved positions at once.

90 We study the blockwise schedule used by LLaDA [2, 27]. Blocks finish from left to right, while a
 91 confidence rule decides which masks inside the active block commit at each step. We write B_b for
 92 the active block and F_b for the completed prefix before it. Shibboleth keeps both the block order and
 93 this confidence rule.

94 3.2 Problem Setup

95 The goal is to distinguish text generated with a secret document key from text generated independently
 96 of that key, at a false-positive rate fixed in advance. The watermark may change which token is
 97 sampled, but it should not replace the model’s decoder or its schedule.

98 **What the two parties retain.** The embedder has the model, prompt, secret key K , and target bits.
 99 The verifier later has the completed sequence, the same checkpoint, K , and public parameters—but
 100 no generation trace, stored logits, random seed, or record of rejected samples. A nonce derives a fresh
 101 document key $K_d = \text{HMAC}(K, \text{nonce}_d)$ [5]. Separate hash domains use K_d for probe construction,
 102 carrier assignment, and the requested response directions.

103 **Guarantee and scope.** For any text produced independently of K_d , including output from another
 104 sampler, a detector run at level α fires with probability at most α . This guarantee is finite-sample and
 105 permits the method to choose a data-dependent number of carriers. It does not promise detection
 106 after editing, key disclosure, adaptive detector queries [16], or replay with a numerically different
 107 checkpoint.

108 **Why generation-time access is needed.** One might instead edit a finished text. Across 36 post-hoc
 109 substitution settings, however, the median selected position has only one candidate within a 3-nat
 110 likelihood budget, and the best setting below a $1.35\times$ PPL budget reaches 0.10 TPR at 1% FPR.
 111 Shibboleth therefore writes while the model’s live distribution is available.

112 The design follows three requirements:

- 113 1. **Replayability.** From the completed sequence, the verifier must reconstruct the same masked
 114 context that the embedder used.
- 115 2. **Blind selection.** Carrier positions must be fixed before the key reveals which response direction
 116 counts as a match.
- 117 3. **Schedule preservation.** Blocks and within-block commit order must remain those of the refer-
 118 ence decoder.

119 Replayability makes a generation-time response measurable later; blind selection keeps data-
 120 dependent carrier discovery from changing the stated false-positive rate. The next three subsections
 121 implement these requirements.




Figure 2: **One block of generation and verification.** Two hidden-prefix views separate candidates by response direction. Carrier selection occurs before the key supplies its orientation; generation commits the first matching draw. Verification rebuilds the same masked state and returns an exact binomial p -value.




Figure 3: **Selection precedes the keyed orientation.** The probe pair divides candidates by response direction. Position i is selected from the two masses alone; only then does K provide ϵ_i . For text produced independently of K , the two signs agree with probability $1/2$. Token labels and bar widths are schematic.

3.3 Probing the Model and Selecting Carriers

For each unfinished block, Shibboleth asks the model the same question under several hidden-prefix views: how likely is each candidate token when this part of the earlier context is missing? A candidate’s response is the change between two such views. A useful carrier is a position where the model still assigns probability mass to candidates with both response directions (Figures 2 and 3).

Formally, Shibboleth constructs L pseudorandom probe patterns $D_{b,\ell} \subset F_b$, where each pattern specifies prefix positions to hide. It also masks the current block B_b and every later position. One model evaluation for pattern ℓ then gives the log-probability of vocabulary item v at position i :

$$\lambda_{i,\ell}(v) = \log p_\theta(v \mid F_b \setminus D_{b,\ell}; B_b \text{ and suffix masked}). \quad (1)$$

The block and suffix are hidden because neither was available when the probe was created. The verifier can therefore reconstruct the same state from the completed sequence instead of seeing later tokens that would alter the probabilities. Probing a block costs $L + 1$ sequence evaluations: L hidden-prefix views and one view with no prefix positions hidden.

For a pair (ℓ, ℓ') , define the response $g_i^{\ell,\ell'}(z) = \lambda_{i,\ell'}(z) - \lambda_{i,\ell}(z)$. Let p_i^{base} be the block-masked distribution with no hidden prefix. The probability mass on either response direction, and a symmetric score for the pair, are

$$q_{i,\pm}^{\ell,\ell'} = \sum_z p_i^{\text{base}}(z) \mathbf{1}[\pm g_i^{\ell,\ell'}(z) > 0], \quad S_i^{\ell,\ell'} = 2 \min(q_{i,+}^{\ell,\ell'}, q_{i,-}^{\ell,\ell'}). \quad (2)$$

The score is high only when both directions remain available. At each position, Shibboleth keeps the pair with the largest S_i and selects the position when $S_i > s_{\min}$, subject to an optional per-block cap. Choosing the best pair raises the measured mean score from 0.28 for a fixed near/far pair to 0.45.

Crucially, swapping ℓ and ℓ' leaves S_i unchanged while reversing every response sign. Carrier selection therefore reveals nothing about which sign will later be rewarded. Figure 3 makes this ordering explicit.

After selecting the carrier, a separate hash domain supplies $\epsilon_i \in \{-1, +1\}$. A target $w_i \in \{-1, +1\}$ represents the provenance or payload bit. Candidate z counts as a match when

$$m_i(z) = \mathbf{1}[w_i \epsilon_i g_i(z) > 0]. \quad (3)$$

If $g_i(z) = 0$, a second domain-separated bit assigns the outcome. Such a tie cannot be steered, but it remains an independent fair outcome for detection.

3.4 Writing the Requested Answer

The reference decoder still chooses the next position and supplies its live distribution p_i . Unselected positions are sampled once as usual. At a selected carrier, Shibboleth accepts a candidate only if it has the requested response and is not too far below the most likely token:

$$A_i = \left\{ z : w_i \epsilon_i g_i(z) > 0 \text{ and } p_i(z) \geq \kappa \max_v p_i(v) \right\}. \quad (4)$$

It draws from the same p_i up to R times and commits the first member of A_i . If every proposal misses, it takes one fresh unchecked draw. The probability row is never recomputed, so retries add no model evaluation and do not alter the commit order.

Let $a_i = p_i(A_i)$ be the mass that passes both conditions, and let $q_i = \sum_z p_i(z) m_i(z)$ be the chance that an unchecked draw scores as a match. The committed token comes from $p_i(\cdot \mid A_i)$ with probability $1 - (1 - a_i)^R$ and from p_i otherwise. Its match probability is therefore

$$\Pr[m_i = 1] = 1 - (1 - a_i)^R (1 - q_i). \quad (5)$$

With no plausibility floor and no ties, $a_i = q_i$, giving $1 - (1 - q_i)^{R+1}$. Increasing R strengthens detection but changes the sampling distribution; κ limits how unlikely a guided proposal may be relative to the live argmax. It does not constrain the unchecked fallback or guarantee semantic quality, which we measure separately.

3.5 Watermark Detection

The verifier re-masks each completed block and suffix, reconstructs the probe pairs and carriers, and evaluates $m_i(y_i)$. It needs no rejected samples or generation trace. For provenance detection, the selected carriers form a pool P with $w_i = +1$; a payload implementation would keep a separate pool for each message bit. The document statistic and its one-sided p -value are

$$T = \sum_{i \in P} m_i, \quad p_{\text{det}} = \Pr[\text{Binomial}(|P|, \frac{1}{2}) \geq T]. \quad (6)$$

Proposition 1 (Exact false-positive control). Assume the hash domains behave as independent random functions. For any text written without the document key, condition on the selected carriers, probe pairs, responses, and target bits. The orientation bits remain independent fair coins. Hence the m_i are independent Bernoulli($\frac{1}{2}$) variables, $T \sim \text{Binomial}(|P|, \frac{1}{2})$, and thresholding Eq. (6) at α gives false-positive rate at most α for every document length and carrier count.

Why adaptive carrier selection is allowed. The text and model may determine how many carriers exist, which probe pair is chosen, and how large each response is. All of those choices are fixed before the independent orientation bit is used. Flipping ϵ_i then flips the match whenever $g_i(y_i) \neq 0$; the tie rule supplies a fresh fair bit otherwise. The detector therefore needs neither an asymptotic approximation nor a calibration corpus. The guarantee still assumes text produced independently of the key and domain separation among the key’s several roles.

4 Experiments

Models and data. We evaluate GSAI-ML/LLaDA-8B-Instruct [27] and Dream-7B-Instruct [38] in fp16 on one TITAN RTX. Detectability




Figure 4: **Document-level detection.** Replayed carriers (blue outlines) determine T ; the highlighted binomial tail gives the decision.

Method	Setting	PPL ratio ↓	TPR at FPR ↑		Sequence evaluations ↓	
			1%	0.1%	Generation	Detection
<i>LLaDA-8B-Instruct</i>						
KGW [18]	$\delta = 1$	1.22×	0.90	0.83	512	0
dgMARK [14]	$k = 1$	1.45×	0.83	0.80	512	0
dgMARK [14]	+3-beam	1.21×	0.93	0.93	~2048	0
Shibboleth	$R = 1$	0.97×	0.43	0.20	656	144
Shibboleth	$R = 8, \kappa = 0.1$	0.87×	0.70	0.60	656	144
Shibboleth	$R = 16, \kappa = 0.05$	0.74 ×	0.80	0.60	656	144
<i>Dream-7B-Instruct</i>						
KGW [18]	$\delta = 1$	1.84×	0.57	0.37	512	0
dgMARK [14]	$k = 1$	1.75×	0.37	0.27	512	0
dgMARK [14]	+3-beam	0.71 ×	0.30	0.30	~2048	0
Shibboleth	$R = 1$	1.34×	0.40	0.27	656	144
Shibboleth	$R = 8, \kappa = 0.1$	0.79×	0.70	0.67	656	144
Shibboleth	$R = 16, \kappa = 0.05$	1.09×	0.90	0.90	656	144

Table 1: **Detectability and cost at 512 tokens.** PPL ratio is measured against each method’s own unwatermarked run and should be interpreted within that method. Shibboleth uses $n = 30$ documents ($n = 20$ for $R = 16$); Dream PPL uses 8–13 valid pairs per arm. Settings chosen on LLaDA are transferred to Dream without retuning. δ is KGW’s logit bias and k is dgMARK’s candidate count.

Method	MMLU Acc ↑	GSM8K Acc ↑	HumanEval Pass@1 ↑
<i>LLaDA-8B-Instruct</i>			
KGW [18]	0.530→0.535	0.640→0.740	0.329→0.293
dgMARK [14]	0.605→0.530	0.740→0.660	—
dgMARK +3-beam	0.605→0.530	0.740→0.660	—
Shibboleth	0.595→0.575	0.640→0.700	0.335→0.384
<i>Dream-7B-Instruct</i>			
KGW	0.640→0.625	0.700→0.700	—
dgMARK	0.740→0.570	0.840→0.560	—
dgMARK +3-beam	0.740→0.645	0.840→0.560	—
Shibboleth	0.695→0.710	0.660→0.680	—

Table 2: **Task accuracy before and after watermarking.** Each cell is unwatermarked → watermarked under that method’s decoder. MMLU uses 200 questions, GSM8K 50 problems, and HumanEval 164 executed tasks. Approximate standard errors are 0.035, 0.07, and 0.037, respectively; — denotes an unmeasured pair. Shibboleth uses $R = 16, \kappa = 0.05$.

uses C4 realnewslike prompts [32], truncated to 300 characters, and 512-token continuations divided into 32-token blocks. Dream predicts position i at raw-logit index $i - 1$; shifting the wrapper by one restores the convention used for both generation and replay.

Decoders and settings. Shibboleth and dgMARK use each model’s confidence schedule with multinomial sampling at temperature 0.8; dgMARK additionally uses top- $k = 3$. KGW requires strictly left-to-right decoding. Shibboleth uses $L = 8$ probe patterns, threshold $s_{\min} = 0.5$, retry budget R , and optional plausibility floor κ . We evaluate provenance ($w_i = +1$), not payload recovery. With 16 blocks and $L + 1 = 9$ probe views, verification costs 144 sequence evaluations and generation costs $512 + 144 = 656$.

Metrics. Detection is TPR at the analytic FPR thresholds in Eq. (6). Perplexity is scored by GPT-2-large [31] as $\exp(\text{mean}(\text{NLL}_{\text{wm}} - \text{NLL}_{\text{unwm}}))$. A PPL ratio therefore compares a method with its own unwatermarked decoder; it is not a semantic-quality score and is only indicative across different decoders. The baselines are KGW [18] and dgMARK [14].

4.1 Main Results and Analysis

Retries trade sampling distortion for detection. On LLaDA, $R = 1$ reaches 0.43 TPR at 1% FPR. On 30 held-out prompts, $R = 2, 4, 8$ increase TPR to 0.73, 0.83, 0.87, consistent with Eq. (5). Retries add no model evaluation, but without a floor the PPL ratio rises from $1.14\times$ at $R = 2$ to $1.61\times$ at $R = 8$. The floor removes low-probability guided draws: $(R, \kappa) = (8, 0.1)$ gives 0.70 TPR

at $0.87\times$ PPL, and $(16, 0.05)$ gives 0.80 at $0.74\times$. Ratios below one mean only that GPT-2-large prefers the retained high-probability tokens.

The same settings behave differently across checkpoints. We transfer all LLaDA settings to Dream without retuning. KGW falls to 0.57 TPR at 1% FPR and produces 84% duplicate bigrams under forced left-to-right decoding. Direct and three-beam dgMARK reach 0.37 and 0.30 TPR. Shibboleth at $(16, 0.05)$ reaches 0.90 at both reported FPR levels. Its PPL ratio is $0.79\times$ at $(8, 0.1)$ and $1.09\times$ at $(16, 0.05)$, based on $8\text{--}13$ valid pairs per arm. Detection transfers here; the PPL-minimising setting does not.

Task accuracy exposes decoder-specific costs. Table 2 shows each method’s own unwatermarked and watermarked scores side by side; the full grid, including greedy-decoding cells, is Table ?? in Appendix ?. For Shibboleth, changes are at most 0.06 across the measured MMLU, GSM8K, and HumanEval cells. For direct dgMARK on Dream, MMLU changes from 0.740 to 0.570 and GSM8K from 0.840 to 0.560 . These task outputs are also shorter: on LLaDA GSM8K, Shibboleth and direct dgMARK reach 0.12 TPR at 1% FPR, three-beam dgMARK 0.44 , and KGW zero. The 512 -token detection rates therefore do not describe 256 -token answers.

Model-response detection costs model calls. At 512 tokens, three-beam dgMARK reaches 0.93 TPR at both 1% and 0.1% FPR, direct dgMARK $0.83/0.80$, KGW $0.90/0.83$, and Shibboleth $0.80/0.60$. The token-local methods need no model evaluation during detection; Shibboleth needs 144 . KGW’s left-to-right decoder also repeats heavily: before bigram deduplication, its unwatermarked run fires on 37% of documents, and after deduplication two of thirty documents still cross the 0.1% threshold. With cohorts of $20\text{--}30$ documents, TPR differences near 0.1 remain imprecise.

4.2 Calibration and Robustness

Observed false positives remain below the requested levels. Across 20 keys and 120 unwatermarked LLaDA generations per key (Table ??, Appendix ??), mean FPR is 0.0121 , 0.0037 , and 0.0000 at nominal 0.10 , 0.05 , and 0.01 . The largest per-key rates are 0.0583 , 0.0250 , and 0.0000 . Every Dream unwatermarked cohort has zero false positives at 1% .

Editing mainly breaks prefix synchronisation. At $R = 1$, same-model argmax re-denoising or deleting a random 10% of positions reduces TPR at 1% FPR from 0.35 to 0.05 on 20 documents (Table ??, Appendix ??). At $R = 16$, $\kappa = 0.05$ and 1024 tokens ($n = 16$), a random 5% edit leaves TPR at 0.88 , while a random 10% edit reduces it to 0.50 ; false-positive control remains valid on identically edited unwatermarked text. In contrast, re-denoising one contiguous span covering 10% or 30% of positions leaves TPR at 0.88 , because carriers before the span are unchanged.

Neither tested aggregation change improves the random-edit result. A Bonferroni-combined per-block test gives 0.38 TPR versus 0.50 for document-wide counting under random 10% editing, and 0.69 versus 0.88 without editing. A 128 -token conditioning window gives 0.44 versus 0.50 under random 10% editing. Copy-paste-style sparse survival was not evaluated.

5 Discussion

When is model-response watermarking useful? Shibboleth is most appropriate when the production decoder should remain unchanged and the verifier can retain the checkpoint. It does not replace a token-local detector: the latter is cheaper and easier to run at scale. A practical deployment could use a token-local test for routine checks and reserve checkpoint replay for documents that require a model-specific provenance claim.

Checkpoint binding is both the signal and the cost. A successful Shibboleth test says that the text agrees with keyed responses of a particular checkpoint, not merely that its token identities follow a keyed pattern. This is the method’s distinguishing property, but it also makes the checkpoint identity, numerical precision, and replay implementation part of the verification record. Transfer across numerically different copies is not established by our experiments.

250 **Editing reveals where synchronisation lives.** Dispersed edits alter many later prefixes and there-
251 fore many reconstructed probes; a contiguous span leaves all earlier carriers intact. The measured
252 asymmetry suggests that content-anchored probes, rather than absolute prefix positions, are the most
253 direct route to editing tolerance. The block-local and windowed variants tested here did not improve
254 detection after random edits, so this remains future work rather than a demonstrated remedy.

255 **Interpreting quality.** A perplexity ratio below one does not by itself imply better text, and the
256 task cohorts are too small to support claims about modest accuracy changes. The relevant result
257 is narrower: under the measured settings, Shibboleth’s task score changes are small relative to the
258 degradation observed for direct dgMARK on Dream, while its detector requires substantially more
259 computation.

260 6 Conclusion

261 Shibboleth shows that a masked DLM can participate in authenticating its own output. The watermark
262 is carried by reproducible changes in model probability under hidden context, rather than by a keyed
263 token subset. Its central design choice is temporal: select a carrier from a symmetric response score
264 first, then reveal the keyed orientation. This ordering gives an exact binomial test even though the
265 text and model determine how many carriers are available.

266 The result complements token-local watermarking rather than replacing it. Shibboleth preserves the
267 deployed decoding schedule and transfers strongly from LLaDA to Dream, but its verifier must rerun
268 the same checkpoint 144 times per 512-token document. Token-local detectors remain cheaper and,
269 on LLaDA, reach higher TPR.

270 **Limitations.** The false-positive guarantee does not imply detection after editing. Random 10%
271 re-denoising reduces TPR from 0.35 to 0.05 at $R = 1$ and from 0.88 to 0.50 at $R = 16$, although
272 one contiguous edited span leaves the measured TPR unchanged. Replay also requires the same
273 checkpoint and numerical precision. Evaluation covers two masked DLMs, one prompt corpus,
274 cohorts of 16–50 documents, and provenance only; payload recovery is not measured.

References

- [1] Scott Aaronson and Hendrik Kirchner. Watermarking GPT outputs, 2023. <https://www.scottaaronson.com/talks/watermark.ppt>.
- [2] Marianne Arriola, Subham Sekhar Sahoo, Aaron Gokaslan, Zhihan Yang, Zhixuan Qi, Jiaqi Han, Justin T. Chiu, and Volodymyr Kuleshov. Block diffusion: Interpolating between autoregressive and diffusion language models. In *ICLR*, 2025.
- [3] Jacob Austin, Daniel D. Johnson, Jonathan Ho, Daniel Tarlow, and Rianne van den Berg. Structured denoising diffusion models in discrete state-spaces. In *NeurIPS*, 2021.
- [4] Arka Bagchi, Akhil Bhimaraju, Moulik Choraria, Daniel Alabi, and Lav R. Varshney. Watermarking discrete diffusion language models. *arXiv preprint arXiv:2511.02083*, 2025.
- [5] Mihir Bellare, Ran Canetti, and Hugo Krawczyk. Keying hash functions for message authentication. In *CRYPTO*, 1996.
- [6] Heli Ben-Hamu, Itai Gat, Daniel Severo, Niklas Nolte, and Brian Karrer. Accelerated sampling from masked diffusion models via entropy bounded unmasking. In *NeurIPS*, 2025.
- [7] Adam Block, Alexander Rakhlin, and Ayush Sekhari. GaussMark: A practical approach for structural watermarking of language models. In *ICML*, 2025.
- [8] Ruibo Chen, Yihan Wu, Yanshuo Chen, Chenxi Liu, Junfeng Guo, and Heng Huang. A watermark for order-agnostic language models. In *ICLR*, 2025.
- [9] Miranda Christ, Sam Gunn, and Or Zamir. Undetectable watermarks for language models. In *COLT*, 2024.
- [10] Pierre Fernandez, Antoine Chaffin, Karim Tit, Vivien Chappelier, and Teddy Furon. Three bricks to consolidate watermarks for large language models. In *2023 IEEE International Workshop on Information Forensics and Security (WIFS)*, 2023.
- [11] Jiayi Fu, Xuandong Zhao, Ruihan Yang, Yuansen Zhang, Jiangjie Chen, and Yanghua Xiao. GumbelSoft: Diversified language model watermarking via the GumbelMax-Trick. In *ACL*, 2024.
- [12] Thibaud Gloaguen, Robin Staab, Nikola Jovanović, and Martin Vechev. Watermarking diffusion language models. In *ICLR*, 2026. arXiv:2509.24368.
- [13] Jonathan Ho, Ajay Jain, and Pieter Abbeel. Denoising diffusion probabilistic models. In *NeurIPS*, 2020.
- [14] Pyo Min Hong and Albert No. dgMARK: Decoding-Guided Watermarking for Diffusion Language Models. In *ICML*, 2026. arXiv:2601.22985.
- [15] Zhengmian Hu, Lichang Chen, Xidong Wu, Yihan Wu, Hongyang Zhang, and Heng Huang. Unbiased watermark for large language models. In *The Twelfth International Conference on Learning Representations (ICLR)*, 2024.
- [16] Nikola Jovanović, Robin Staab, and Martin Vechev. Watermark stealing in large language models. In *ICML*, 2024.
- [17] Jaeyeon Kim, Kulin Shah, Vasilis Kontonis, Sham M. Kakade, and Sitan Chen. Train for the worst, plan for the best: Understanding token ordering in masked diffusions. In *ICML*, 2025.
- [18] John Kirchenbauer, Jonas Geiping, Yuxin Wen, Jonathan Katz, Ian Miers, and Tom Goldstein. A watermark for large language models. In *ICML*, 2023.
- [19] John Kirchenbauer, Jonas Geiping, Yuxin Wen, Manli Shu, Khalid Saifullah, Kezhi Kong, Kasun Fernando, Aniruddha Saha, Micah Goldblum, and Tom Goldstein. On the reliability of watermarks for large language models. In *ICLR*, 2024.

- [20] Kalpesh Krishna, Yixiao Song, Marzena Karpinska, John F. Wieting, and Mohit Iyyer. Paraphrasing evades detectors of AI-generated text, but retrieval is an effective defense. In *NeurIPS*, 2023.
- [21] Rohith Kuditipudi, John Thickstun, Tatsunori Hashimoto, and Percy Liang. Robust distortion-free watermarks for language models. *TMLR*, 2024.
- [22] Inception Labs, Samar Khanna, Siddhant Kharbanda, Shufan Li, Harshit Varma, Eric Wang, Sarah Birnbaum, Ziyang Luo, Yanis Miraoui, and Akash Palrecha. Mercury: Ultra-fast language models based on diffusion. *arXiv preprint arXiv:2506.17298*, 2025.
- [23] Chia-Ming Lee, Shao-Kai Liu, Ming-Ching Chang, Xin Li, Yu-Lun Liu, and Chih-Chung Hsu. Commit locally, exit globally: Coordinating adaptive sampling and early exit in diffusion language models, 2026. URL <https://arxiv.org/abs/2607.28166>.
- [24] Xiang Li, Feng Ruan, Huiyuan Wang, Qi Long, and Weijie J. Su. Robust detection of watermarks for large language models under human edits. *Journal of the Royal Statistical Society: Series B*, 2025.
- [25] Yuqing Liang, Jiancheng Xiao, Wensheng Gan, and Philip S. Yu. Watermarking techniques for large language models: A survey. *Artificial Intelligence Review*, 2026.
- [26] Aaron Lou, Chenlin Meng, and Stefano Ermon. Discrete diffusion modeling by estimating the ratios of the data distribution. In *ICML*, 2024.
- [27] Shen Nie, Fengqi Zhu, Zebin You, Xiaolu Zhang, Jingyang Ou, Jun Hu, Jun Zhou, Yankai Lin, Ji-Rong Wen, and Chongxuan Li. Large language diffusion models. In *NeurIPS*, 2025.
- [28] Jingyang Ou, Shen Nie, Kaiwen Xue, Fengqi Zhu, Jiacheng Sun, Zhenguo Li, and Chongxuan Li. Your absorbing discrete diffusion secretly models the conditional distributions of clean data. In *ICLR*, 2025.
- [29] Fabien A. P. Petitcolas, Ross J. Anderson, and Markus G. Kuhn. Information hiding — a survey. In *Proceedings of the IEEE*, 1999.
- [30] Ohad Raban, Ethan Fetaya, and Gal Chechik. LR-DWM: Efficient watermarking for diffusion language models. *arXiv preprint arXiv:2601.12376*, 2026.
- [31] Alec Radford, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever. Language models are unsupervised multitask learners. Technical report, OpenAI, 2019.
- [32] Colin Raffel, Noam Shazeer, Adam Roberts, Katherine Lee, Sharan Narang, Michael Matena, Yanqi Zhou, Wei Li, and Peter J. Liu. Exploring the limits of transfer learning with a unified text-to-text transformer. *Journal of Machine Learning Research*, 21(140):1–67, 2020.
- [33] Subham Sekhar Sahoo, Marianne Arriola, Aaron Gokaslan, Edgar Marroquin Marroquin, Alexander M. Rush, Yair Schiff, Justin T. Chiu, and Volodymyr Kuleshov. Simple and effective masked diffusion language models. In *NeurIPS*, 2024.
- [34] Jiaxin Shi, Kehang Han, Zhe Wang, Arnaud Doucet, and Michalis Titsias. Simplified and generalized masked diffusion for discrete data. In *NeurIPS*, 2024.
- [35] Qingyan Wei, Yaojie Zhang, Zhiyuan Liu, Pengfei Zeng, Yu Wang, Biao Qi, Dong Liu, and Linfeng Zhang. Accelerating diffusion large language models with SlowFast sampling: The three golden principles. In *ICLR*, 2026.
- [36] Chengyue Wu, Hao Zhang, Shuchen Xue, Zhijian Liu, Shizhe Diao, Ligeng Zhu, Ping Luo, Song Han, and Enze Xie. Fast-dLLM: Training-free acceleration of diffusion LLM by enabling KV cache and parallel decoding. In *ICLR*, 2026.
- [37] Linyu Wu, Linhao Zhong, Wenjie Qu, Yifan Li, Yue Liu, Shengfang Zhai, Chunhua Shen, and Jiaheng Zhang. DMark: Order-agnostic watermarking for diffusion large language models. *arXiv preprint arXiv:2510.02902*, 2025.

- 365 [38] Jiacheng Ye, Zhihui Xie, Lin Zheng, Jiahui Gao, Zirui Wu, Xin Jiang, Zhenguo Li, and Lingpeng
366 Kong. Dream 7B: Diffusion large language models. *arXiv preprint arXiv:2508.15487*, 2025.
- 367 [39] Xuandong Zhao, Yu-Xiang Wang, and Lei Li. Protecting language generation models via
368 invisible watermarking. In *ICML*, 2023.
- 369 [40] Xuandong Zhao, Prabhanjan V. Ananth, Lei Li, and Yu-Xiang Wang. Provable robust water-
370 marking for AI-generated text. In *ICLR*, 2024.
- 371 [41] Fengqi Zhu, Rongzhen Wang, Shen Nie, Xiaolu Zhang, Chunwei Wu, Jun Hu, Jun Zhou, Jianfei
372 Chen, Yankai Lin, Ji-Rong Wen, and Chongxuan Li. LLaDA 1.5: Variance-reduced preference
373 optimization for large language diffusion models. *arXiv preprint arXiv:2505.19223*, 2025.